

qicert: Certified Compression of Vision–Language–Action Models

with Exact Lipschitz Bounds and a Refusing Runtime Guard

Phase-1 Concept Proposal — 2026 Global Quantum + AI Challenge

Problem statement: *Quantum-Enhanced Vision-Language-Action Models in Autonomous Driving and Robotics Applications* (Volkswagen Group)

Code: <https://github.com/ZakLr/qicert> · Technical report: [docs/technical-report-v2.pdf](#) (7 pages, in repo)

1. Problem Framing

Vision–language–action models (VLAMs) for driving and robotics are 7B–10B parameter systems that must eventually run on embedded controllers under functional-safety regimes (ISO 26262, IEC 62061). The challenge statement names the compound constraint precisely: shrink the model *and* retain verifiable behavior on long-tail inputs. Classical compression solves the first half only. Quantization, pruning, and distillation reduce footprint, but none of them yields a statement of the form “under any input perturbation of size δ , the action cannot leave this set” — which is the artifact a safety engineer would need before certifying deployment on long-tail sensor distributions.

Quantum or quantum-inspired advantage enters where structure, not scale, does the work. Our QI component is tensor-network structure (tensor-train decomposition, the MPS/TTNS family the statement highlights). Its decisive property is not size reduction — at the ratios we tested it loses the accuracy contest to INT8 and we report that plainly — but that the compressed weights *carry an exact per-layer Lipschitz constant* computable from the compressed cores alone, with no reference to data. Quantized weights admit no comparable exact bound, because quantization error is data-dependent. Structure is therefore the classical-side object that makes an exact deployment certificate computable, chainable, and enforceable at runtime — and it is structurally unavailable to the accepted baseline. This is the specific, narrow sense in which the quantum-inspired method provides the credible advantage the guidelines ask for; we deliberately claim nothing broader.

2. Technical Approach

Paradigm: quantum-inspired (tensor-network), classical pipeline, no quantum hardware required or assumed. The statement’s simulation clause covers us explicitly; we state resources honestly in Section 3.

The pipeline wraps an existing VLA at four points. A fine-tuning stage establishes the reference policy and freezes a hash-verified evaluation split. The compression stage decomposes each linear layer of the language backbone into a tensor-train network (TT-SVD); a closed-form low-rank residual folds the truncation loss back into the operator so that accuracy survives compression as well as the budget allows, with the certificate cost known exactly (triangle inequality on the operator norm). The certificate stage chains the exact per-layer bounds into a model-level perturbation ball and signs the artifact into a hash manifest. The deployment stage is a runtime guard that verifies the manifest at load time and, at every inference step, refuses any action outside the certified ball.

Three measured components complete the stack. First, a commuting-Pauli interaction compiler evaluates cross-modal interaction operators exactly over pruned families; its pruning predictions track measurement at median $r = 0.975$ over twelve independent operator tables, with compiler identity verified to 2.3×10^{-14} relative error. Second, a rare-event estimation layer races three estimators on a known-ground-truth problem ($p^* = 1.37 \times 10^{-4}$): an iterated quantum amplitude estimation schedule (simulated; the estimator sees only measurement counts, as it would on hardware) reaches naive-Monte-Carlo interval width with $\sim 1,587\times$ fewer queries, and all arms cover the truth at budget. Third, a conformal layer calibrates violation thresholds whose empirical coverage lands inside the finite-sample band at all three tested levels ($\alpha = 0.01/0.05/0.10$).

Proof status. The layer bound $\|W\|_2 \leq \prod_k \|G^{(k)}\|_3$ is exact algebra on the TT factorization; the certificate chain and the guard’s soundness/completeness properties are machine-checked by SMT falsification (Z3 finds no counterexample to accepting an out-of-ball action). A guard demo transcript shows: clean load SERVE, out-of-ball action REFUSE, tampered-certificate REFUSE, NaN-bound REFUSE, missing-manifest REFUSE.

Working prototype. This is not a design on paper. The repository contains 80 passing tests, a benchmark harness, and measured results on a fine-tuned 0.5B VLA (LIBERO-spatial, 114,497 scored action tokens, hash-frozen split): reference accuracy 0.4468; calibrated INT8 at $1.997\times$ compression, 0.4466 (the honest classical comparator); tensor-compressed routes at $2.5\text{--}4.1\times$ with sound certificates on every one of 168 compressed layers across every configuration tested. The accuracy of the repaired tensor route under the pre-registered bar is the open question Section 5 turns into the Phase-2 validation plan; the negative result at the tested configurations is reported in full in the technical report rather than hidden.

3. Feasibility and Resource Requirements

Feasibility is demonstrated, not projected: every component above has already run end to end on consumer hardware. The full experimental program to date consumed roughly 15 GPU-hours on a single laptop-class GPU (NVIDIA RTX 5060, 8 GB), including fine-tuning (1,200 steps), all compression arms, full-split evaluations, the estimator fixtures, and latency profiling. Peak training memory was 4.8 GiB. No proprietary data, cloud credits, or quantum hardware are involved at any stage; LIBERO is the statement’s own recommended robotics dataset; simulation environments and versions are pinned in the repository. Phase-2 scale-up (below) is sized for the same single-GPU discipline, with multi-seed replication the main added cost. The honest constraint we declare: tensor-network contraction on a CPU edge profile costs $3\text{--}15\times$ the dense fp16 kernel time per layer at these shapes — within the 100 ms budget on measured paths, but a real margin cost that Phase 2 must optimize rather than wish away.

4. Expected Impact

For the enterprise use case, the deliverable changes what deployment approval means. A fleet operator today ships a compressed model on the strength of benchmark scores; with qicert-style artifacts they would ship a signed certificate, a guard that refuses out-of-ball actions, and an audited statistical tail — the difference between “we tested it” and “it is bounded, and the machine enforces the bound.” The design tool built on our degradation measurements points the same direction: predict the largest safe compression *before* committing to full evaluation, turning post-hoc luck into engineering process.

Quantitative targets, anchored to what already runs: the reference VLA at 0.4468 accuracy and the INT8 comparator at $1.997\times$ define the bar the tensor route must clear; certificates already hold at $4.1\times$ (sound 168/168); the estimator layer already delivers the $\sim 1,587\times$ query reduction. A successful Phase-2 demonstrates a certified configuration that is simultaneously (a) at or beyond $2\times$ compression, (b) within the pre-registered accuracy band (reference -0.05), and (c) certificate-sound with guard enforcement — or, failing that, the certified sub- $2\times$ regime with a validated predictor that says exactly where the accuracy wall sits, which is itself deployable information the quantized baseline cannot produce.

5. Validation Plan

The Phase-2 sprint validates three claims, each with metrics fixed in advance (pre-registered in the repository before the runs they judge) — this is the strongest differentiator we can offer: the gates are committed, mechanical, and auditable.

Capability claim (the live experiment). Configuration search over the repaired tensor-compression route — activation-weighted residual fitting, mixed quantized/tensor allocation, and sub- $2\times$ targets where INT8 is byte-capped — with each promising configuration confirmed on the full 6,496-batch frozen protocol. Metrics: compression ratio, token-level action accuracy, pre-registered gate (ratio $\geq 2\times$ AND accuracy ≥ 0.3968), mean $\pm$ standard deviation over 3 independent seeds (seeds 1–2 are queued at submission time; seed 0 measured).

Safety claim. Metrics: certificate soundness (currently 168/168 across all configurations), guard refusal behavior (currently 1.0 on all four tampering scenarios plus in/out-of-ball cases), estimator coverage of known ground truth (currently 1.0 at budget for all arms), conformal coverage within finite-sample bands (currently true at all three levels).

Design-tool claim. The degradation predictor: fit weight-space features to measured agreement, report leave-one-out R^2 and predicted-vs-measured maximum safe ratio. Success: predicted safe ratio within a stated tolerance of measurement on held-out points.

Integration milestones for the sprint, in order: multi-seed replication; degradation-curve measurement and predictor validation; mixed-allocation search to the gate; edge-profile latency optimization for the tensor contraction (the known margin cost); and a hardware-pathway write-up estimating qubit count, circuit depth, and shot budget for the amplitude estimation component, which the challenge lists as a secondary objective and for which our simulation fixtures already define the interface.

6. Hybrid / Cross-Domain Integration

The design is hybrid by construction, and the division of labor is the point. Classical deep-learning infrastructure does the training and the heavy evaluation (PyTorch/CUDA on consumer GPU). The quantum-inspired tensor layer sits inside the model, where its algebraic structure is load-bearing: it is what makes the exact bound computable at all. Quantum algorithmic structure appears as the amplitude-estimation schedule for the statistical tail — executed in simulation against the same measurement-count interface it would have on hardware, so the estimator layer is hardware-swappable without redesign. Formal methods (SMT falsification via Z3, conformance calibration, scenario optimization) audit the whole stack mechanically rather than by argument. Cross-domain transfer is native to the method: the certificate machinery is backbone-agnostic and dataset-agnostic, so the same pipeline applies unchanged to the autonomous-driving track (the latency profile we measure is a CPU edge profile, and the challenge’s 100 ms real-time constraint is the same in both tracks); the robotics-track instantiation here was chosen because its accepted baseline is reproducible on our hardware.

7. Team Capability

The team combines the three disciplines this proposal straddles. The lead (Zaki) built and operates the full pipeline end to end: the fine-tuning harness, the tensor-compression kernel stack with its certificate algebra, the benchmark/ledger infrastructure with pre-registered gates, and the formal-methods integration — evidenced by the public repository’s measured results, 80-test suite, and reproducible artifact chain. Working practice follows the discipline the reviewers will look for: every reported number resolves to a run directory with environment snapshot; evaluation splits are frozen and hashed before experiments; negative results are reported first-class (the technical report’s headline compression finding is an honest NO-GO under the pre-registered bar, with the failure characterized rather than framed). AI research agents were used as engineering leverage under the lead’s direction; every claim is traceable to committed artifacts, and the experiment log records what was run, when, and why for every entry.

Enclosures

- **Public repository:** <https://github.com/ZakLr/qicert> — README quickstart, pinned environment, 80-test suite, benchmark harness, experiment log, and every artifact behind the numbers above.
- **Technical report (7 pages):** `docs/technical-report-v2.pdf` in the repository — method, theory, measured results incl. the negative capability result, ablation isolating the QI component, limitations, and a claim-to-artifact provenance map (Appendix A) with seven artifact-generated figures.
- **Team profile:** submitted via the portal form (names, roles, affiliations, prior quantum/QI and domain experience) per guidelines §4.1.